

HIGH-PERFORMANCE SCIENTIFIC COMPUTING (HPSC) ME522

Instructor: Gaurav Bhutani

SMME, IIT Mandi

Feb-Jun 2023 semester

Venue: to be announced

Online: <https://meet.google.com/xyn-osvu-yy>

Lecture plan

- Course essentials
- Goals and learning outcomes
- Pre-requisites and software requirements
- Short demonstration

Course essentials

Notes

- Class: 4 hours per week (2 sessions).
- Venue to be announced.
- One hr lecture, followed by 1 hour lab / hands-on session. Use your own laptop.
- Teaching supporter: Ayush Sahu (SMME Metch
Help with labs, installing software, obtaining material, announcements
- Announcements: Course mailing list / Google chat group
- Learn from peers
- Moodle – all course slides and link to material
- Codes on Github as we create
- Virtual machines on IIT Mandi Cloud

Evaluation

- Lab/class attendance and random lab viva (20)
- Midsem (30)
 - Written (10)
 - Lab exam (20)
- Endsem (50)
 - Written (20)
 - Lab exam (30)

Goals

HPC

- **High Performance Computing (HPC)** often means heavy-duty computing on clusters or supercomputers with 100s of thousands of cores.
- *“World’s fastest computer”*
- #1. Frontier – HPE Cray (Oak Ridge National Lab, USA)
8.7M cores \approx 1100 Petaflops; 21MW power
- #3. Leonardo – Atos (CINECA, Italy)
1.4M cores \approx 174 Petaflops; 5.6MW
- Param Himalaya – Atos (IIT Mandi, India)
800 Teraflops; 150kW


See <http://top500.org> for current list.



Increasing speed

- Moore's Law: Processor speed doubles every 18 months.
⇒ factor of 1024 in 15 years.
- Going forward: Number of cores doubles every 18 months
- Top: Total computing power of top 500 computers
- Middle: #1 computer
- Bottom: #500 computer

<http://www.top500.org>



Our focus

- Our focus is more modest, but we will cover material that is:
 - Essential to know if you eventually want to work on supercomputers
 - Extremely useful for any scientific computing project, even on a laptop.
- Focus on scientific computing as opposed to other computationally demanding domains, for which somewhat different tools might be best.

Learning outcomes

Efficient processing

- Basic computer architecture, e.g. floating point arithmetic, cache hierarchies
- Using Unix (or Linux)
- Language issues, e.g. compiled vs. interpreted, object oriented, etc.
- Specific languages: Python (for scripting), Fortran 90/95 (for fast processing)
- Parallel computing with OpenMP, MPI
- Not included: GPU-based parallelisation

Good software practices

- Version control using Git, github
- Makefiles
- Debuggers, code testing
- Reproducibility
- Use of high-performance computing (HPC) clusters

Strategy

- So much material, so little time...
- Concentrate on basics, simple motivating examples.
- Get enough hands-on experience to be comfortable experimenting further and learning much more on your own.
- Learn what's out there to help select what's best for your needs. New languages are introduced with time – similar ideas though.
- Teach many things “by example” as we go along.
- You'll be expected to read supplementary notes when they are provided.
- No specific book for the course. Internet search for help will be useful.

Pre-requisites & software requirements

Pre-requisites

- Some programming experience in some language, e.g., Matlab, C.
- You should be comfortable:
 - editing a file containing a program and executing it,
 - using basic structures like loops, if-then-else, input-output,
 - writing subroutines or functions in some language
- You are not expected to know Python or Fortran.
- Some basic knowledge of linear algebra, e.g.:
 - what vectors and matrices are and how to multiply them
 - How to go about solving a linear system of equations
- Comfort level for learning new software.

Software requirements

- You will need access to a computer with a number of things on it. All open-source software.
- Note: Unix is often required for scientific computing.
- Windows: Many tools we'll use can be used with Windows, but learning Unix is part of this class.
- Options:
 - Install everything you'll need on your own computer.
 - Install VirtualBox and use the Virtual Machine (VM) created for this class.
 - Use Amazon Web Services with an Amazon Machine Image (AMI) which will be created for this class. Other cloud computing services are also available.

Know the class

- A short intro by various class members
 - Name and program enrolled in
 - What are your academic interests and what research projects you are working on? For PG students please mention your research group at IIT Mandi.
 - Why interested in this course? What are your expectations from the course?

Short demonstration

Demo code

```
$ cd <folder name>  
$ export HPSC=$PWD  
$ cd $HPSC/lecture1
```

```
$ make plots
```

```
$ display *.png
```


Steady-state heat conduction

- Discretize on an $N \times N$ grid with N^2 unknowns
- Assume temperature is fixed (and known) at each point on boundary.
- At interior points, the steady state value is (approximately) the average of the 4 neighbouring values.

$$(A)_{N^2 \times N^2} \vec{T} = \vec{b}$$

$N^2 \times N^2$

Row-wise ordering




$$\nabla^2 u = 0$$

2-D $\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} = 0$

$$[A]_{N^2 \times N^2} \{x\} = \{b\}$$

$\begin{bmatrix} x_8 \\ x_9 \\ x_{10} \\ x_{11} \\ x_{12} \end{bmatrix}$
 $N^2 \text{ } 120 \times 120$



$$\frac{\partial^2 u}{\partial x^2} = \frac{u_{i+1,j} - 2u_{i,j} + u_{i-1,j}}{(\Delta x)^2}$$

$$u_{i,j} = f(u_{i-1,j}, u_{i+1,j}, u_{i,j-1}, u_{i,j+1})$$

Steady-state heat conduction


$$u_{i,j} = \frac{1}{4} (u_{i-1,j} + u_{i+1,j} + u_{i,j-1} + u_{i,j+1})$$

- Holds for $i, j = 1, 2, \dots, N$ with $u_{0,j}$ known on boundary. Gives a linear system $Au = b$, with N^2 equations N^2 unknowns. Matrix A is $N^2 \times N^2$, for $N = 120$, $N^2 = 14400$.

- Very sparse: each row of matrix A has at most 5 nonzeros. Gaussian elimination is not the best approach.
- Jacobi method (not the best method)

$$u_{(i,j)}^{[k+1]} = \frac{1}{4} (u_{(i-1,j)}^{[k]} + u_{(i+1,j)}^{[k]} + u_{(i,j-1)}^{[k]} + u_{(i,j+1)}^{[k]})$$

Row-wise ordering



Speedup of linear solvers




Fig. 2 Comparison of the contributions of mathematical algorithms and computer hardware.

Source: SIAM Review 43(2001), p. 168.

- Exponential increase in the speed of numerical algorithms.
- Exponential increase in the computer hardware performance with time.

Class virtual machine

- Available on Google Drive -
`https://drive.google.com/drive/folders/1mnMJ0JNqi0pcuFzVZeW8EfoxqlqHWWVG?usp=sharing`
- Username: hpsc; password: me522
- This file is large! About 5 GB compressed. After unzipping, about 10 GB.

References

- Slides are adapted from HPSC, AM483, Uni of Washington by Randall J Leveque
(<https://faculty.washington.edu/rjl/teaching.html>)
Refer to his page for more useful material and notes.