

Outline:

- OpenMP:
- Parallel do loops, reductions

Reading:

- `codes/openmp`

OpenMP parallel do loops

```
!$omp parallel do
do i=1,n
    ! do stuff for each i
enddo
!$omp end parallel do    ! OPTIONAL
```

indicates that the do loop can be done in parallel.

Requires:

- what's done for each value of i is independent of others

- Different values of i can be done in any order.

The iteration variable i is private to the thread: each thread has its own version.

By default, all other variables are shared between threads unless specified otherwise.

Need to be careful that threads use shared variables properly.

OpenMP parallel do loops

This code fills a vector y with function values that take a bit of time to compute:

```
! fragment of codes/openmp/yeval.f90

dx = 1.d0 / (n+1.d0)

!$omp parallel do private(x)
do i=1,n
    x = i*dx
    y(i) = exp(x)*cos(x)*sin(x)*sqrt(5*x+6.d0)
enddo
```

Elapsed time for $n = 10^8$, without OpenMP: about 3.0 sec.

Elapsed time using OpenMP on 2 processors: about 1.9 sec.

Memory stack

Note: Parallel threads use stack and you may need to increase the limit (e.g. on the VM):

```
$ gfortran -fopenmp yeval.f90
```

```
$ ./a.out
```

```
Segmentation fault
```

```
$ ulimit -s
```

```
8192
```

```
$ ulimit -s unlimited
```

```
$ ./a.out
```

```
Using OpenMP with 2 threads
```

```
Filled vector y of length 100000000
```

Memory stack

Note: Parallel threads use stack and you may need to increase the limit (e.g. on the VM):

```
$ gfortran -fopenmp yeval.f90
$ ./a.out
Segmentation fault
```

```
$ ulimit -s
8192
```

```
$ ulimit -s unlimited
```

```
$ ./a.out
Using OpenMP with 2 threads
Filled vector y of length 100000000
```

On Mac, there's a hard limit `ulimit -s hard` or

```
gfortran -fopenmp -Wl,-stack_size -Wl,3f000000 yeval.f90
```

Memory: Heap and Stack

Memory devoted to data for a program is generally split up:

Heap: Dynamically allocated memory — memory allocator looks for free block of memory, keeps track of free list, does garbage collection, etc.

Stack: Block of memory where space is allocated on “top” of the stack as needed and “popped” off the stack when no longer needed. **Last in – first out (LIFO).**

Fast relative to heap allocation.

Natural way to allocate storage for nested subroutine or function calls: If A calls B calls C, then when the variables used by C are popped off the stack, we're back to the variables of B.

Memory: Heap and Stack

Memory devoted to data for a program is generally split up:

Heap: Dynamically allocated memory — memory allocator looks for free block of memory, keeps track of free list, does garbage collection, etc.

Stack: Block of memory where space is allocated on “top” of the stack as needed and “popped” off the stack when no longer needed. **Last in – first out (LIFO).**

Fast relative to heap allocation.

Natural way to allocate storage for nested subroutine or function calls: If A calls B calls C, then when the variables used by C are popped off the stack, we're back to the variables of B.

Private variables for threads also put on stack, popped off when parallel block ends.

OpenMP parallel do loops

This code is **not correct**:

```
!$omp parallel do
do i=1,n
    x = i*dx
    y(i) = exp(x)*cos(x)*sin(x)*sqrt(5*x+6.d0)
enddo
```


OpenMP parallel do loops

This code is **not correct**:

```
!$omp parallel do
do i=1,n
    x = i*dx
    y(i) = exp(x)*cos(x)*sin(x)*sqrt(5*x+6.d0)
enddo
```

By default, `x` is a shared variable.

Might happen that:

- Processor 0 sets `x` properly for one value of `i`,
- Processor 1 sets `x` properly for another value of `i`,
- Processor 0 uses `x` but is now incorrect.

OpenMP parallel do loops

Correct version:

```
!$omp parallel do private(x)
do i=1,n
    x = i*dx
    y(i) = exp(x)*cos(x)*sin(x)*sqrt(5*x+6.d0)
enddo
```

Now each thread has its own version of `x`.

Iteration counter `i` is private by default.

OpenMP parallel do loops

Correct version:

```
!$omp parallel do private(x)
do i=1,n
    x = i*dx
    y(i) = exp(x)*cos(x)*sin(x)*sqrt(5*x+6.d0)
enddo
```

Now each thread has its own version of `x`.

Iteration counter `i` is private by default.

Note that `dx`, `n`, `y` are shared by default. **OK because:**

`dx`, `n` are used but not changed,
`y` is changed, but independently for each `i`

OpenMP parallel do loops

Incorrect code:

```
dx = 1.d0 / (n+1.d0)
!$omp parallel do private(x,dx)
do i=1,n
    x = i*dx
    y(i) = exp(x)*cos(x)*sin(x)*sqrt(5*x+6.d0)
enddo
```

Specifying `dx private` won't work here.

This will create a private variable `dx` for each thread but it will be **uninitialized**.

Will run but give garbage.

OpenMP parallel do loops

Could fix with:

```
dx = 1.d0 / (n+1.d0)
!$omp parallel do firstprivate(dx)
do i=1,n
    x = i*dx
    y(i) = exp(x)*cos(x)*sin(x)*sqrt(5*x+6.d0)
enddo
```

The **firstprivate** clause creates private variables and initializes to the value from the master thread prior to the loop.

OpenMP parallel do loops

Could fix with:

```
dx = 1.d0 / (n+1.d0)
!$omp parallel do firstprivate(dx)
do i=1,n
    x = i*dx
    y(i) = exp(x)*cos(x)*sin(x)*sqrt(5*x+6.d0)
enddo
```

The **firstprivate** clause creates private variables and initializes to the value from the master thread prior to the loop.

There is also a **lastprivate** clause to indicate that the last value computed by a thread (for $i = n$) should be copied to the master thread's copy for continued execution.

OpenMP parallel do loops

```
! from codes/openmp/private1.f90
n = 7
y = 2.d0
!$omp parallel do firstprivate(y) lastprivate(y)
do i=1,n
    y = y + 10.d0
    x(i) = y
    !omp critical
    print *, "i = ",i,"    x(i) = ",x(i)
    !omp end critical
enddo
print *, "At end, y = ",y
```

Run with 2 threads: The 7 values of i will be split up, perhaps

$i = 1, 2, 3, 4$ executed by thread 0,

$i = 5, 6, 7$ executed by thread 1.

Thread 0's private y will be updated 4 times, $2 \rightarrow 12 \rightarrow 22 \rightarrow 32 \rightarrow 42$

Thread 1's private y will be updated 3 times, $2 \rightarrow 12 \rightarrow 22 \rightarrow 32$

OpenMP parallel do loops

```
! from codes/openmp/private1.f90  
n = 7  
y = 2.d0  
!$omp parallel do firstprivate(y) lastprivate(y)  
do i=1,n  
    y = y + 10.d0  
    x(i) = y  
    !omp critical  
    print *, "i = ", i, "    x(i) = ", x(i)  
    !omp end critical  
enddo  
print *, "At end, y = ", y
```

might produce:

i =	1	x(i) =	12.000000000000000
i =	5	x(i) =	12.000000000000000
i =	2	x(i) =	22.000000000000000
i =	6	x(i) =	22.000000000000000
i =	3	x(i) =	32.000000000000000
i =	7	x(i) =	32.000000000000000
i =	4	x(i) =	42.000000000000000
At end, y =			32.000000000000000

Order might be different but final y will be from $i = 7$.

OpenMP parallel do loops — changing default

Default is that loop iterator is private, other variables shared.

Can change this, e.g.

```
!$omp parallel do default(private) shared(x,z) &  
!$omp firstprivate(y) lastprivate(y)  
do i=1,n  
  etc.
```

With this change, only `x` and `z` are shared.

Note continuation character `&` and continuation line.

OpenMP synchronization

```
!$omp parallel do
do i=1,n
    ! do stuff for each i
enddo
!$omp end parallel do    ! OPTIONAL

! master thread continues execution
```

There is an **implicit barrier** at the end of the loop.

The master thread will not continue until all threads have finished with their subset of $1, 2, \dots, n$.

Except if ended by:

```
!$omp end parallel do nowait
```

Conditional clause

Loop overhead may not be worthwhile for short loops.
(Multi-thread version may run slower than sequential)

Can use conditional clause:

```
$omp parallel do if (n > 1000)
do i=1,n
    ! do stuff
enddo
```

If $n \leq 1000$ then no threads are created,
master thread executes loop sequentially.

Nested loops

```
!$omp parallel do private(i)
do j=1,m
  do i=1,n
    a(i,j) = 0.d0
  enddo
enddo
```

The loop on j is split up between threads.

The thread handling $j=1$ does the entire loop on i ,
sets $a(1,1)$, $a(2,1)$, ..., $a(n,1)$.

Nested loops

```
!$omp parallel do private(i)
do j=1,m
  do i=1,n
    a(i,j) = 0.d0
  enddo
enddo
```

The loop on j is split up between threads.

The thread handling $j=1$ does the entire loop on i ,
sets $a(1,1)$, $a(2,1)$, ..., $a(n,1)$.

Note: The loop iterator i must be declared **private**!

j is private by default, i is shared by default.

Nested loops

Which is better? (assume $m \approx n$)

```
!$omp parallel do private(i)
do j=1,m
  do i=1,n
    a(i,j) = 0.d0
  enddo
enddo
```

or

```
do j=1,m
  !$omp parallel do
  do i=1,n
    a(i,j) = 0.d0
  enddo
enddo
```

Nested loops

Which is better? (assume $m \approx n$)

```
!$omp parallel do private(i)
do j=1,m
  do i=1,n
    a(i,j) = 0.d0
  enddo
enddo
```

or

```
do j=1,m
  !$omp parallel do
  do i=1,n
    a(i,j) = 0.d0
  enddo
enddo
```

The first has less overhead: Threads created only once.

The second has more overhead: Threads created m times.

Nested loops

But have to make sure loop can be parallelized!

Incorrect code for replicating first column:

```
!$omp parallel do private(j)
do i=2,n
  do j=1,m
    a(i,j) = a(i-1,j)
  enddo
enddo
```

Corrected: (j 's can be done in any order, i 's cannot)

```
!$omp parallel do private(i)
do j=1,m
  do i=2,n
    a(i,j) = a(i-1,j)
  enddo
enddo
```


Reductions

Incorrect code for computing $\|x\|_1 = \sum_i |x_i|$:

```
norm = 0.d0
!$omp parallel do
do i=1,n
    norm = norm + abs(x(i))
enddo
```

There is a **race condition**: each thread is updating same shared variable `norm`.

Correct code:

```
!$omp parallel do reduction(+: norm)
do i=1,n
    norm = norm + abs(x(i))
enddo
```

A **reduction** reduces an array of numbers to a single value.

Reductions

A more complicated way to do this:

```
norm = 0.d0
!$omp parallel private(mysum) shared(norm)
mysum = 0
!$omp do
do i=1,n
    mysum = mysum + abs(x(i))
enddo

!$omp critical
norm = norm + mysum
!$omp end critical
!$omp end parallel
```

Some other reductions

Can do reductions using $+$, $-$, $*$, \min , \max , .and. , .or. , some others

General form:

```
!$omp parallel do reduction(operator : list)
```

Example with max:

```
y = -1.d300  ! very negative value
!$omp parallel do reduction(max: y)
do i=1,n
    y = max(y,x(i))
enddo
print *, 'max of x = ',y
```

Some other reductions

General form:

```
!$omp parallel do reduction(operator : list)
```

Example with .or.:

```
logical anyzero
```

```
! set x...  
anyzero = .false.
```

```
!$omp parallel do reduction(.or.: anyzero)  
do i=1,n  
    anyzero = anyzero .or. (x(i) == 0.d0)  
enddo  
print *, 'anyzero = ',anyzero
```

Prints T if any $x(i)$ is zero, F otherwise.

Exercise

1. Implement the Monte Carlo Pi code (as given in the midsem exam) in Fortran
2. Parallelise the code using OpenMP