

## Outline:

- Python plus Fortran: f2py
- LAPACK and the BLAS
- Python compilers
- File writing – ASCII vs binary
- Summary

## Sample codes:

- `/codes/f2py`

## f2py — combining Fortran and Python

Often want to use

- Fortran for intensive computations,
- Python to provide nice user interface, plot results, automate a series of runs with different parameters, do convergence tests as grid size is refined, etc.

Can write data files to disk from Fortran, read into Python,  
This is what we've done for plotting in homeworks.

Sometimes nice to call Fortran directly from Python.  
e.g. LAPACK is used under the hood in NumPy.

f2py provides a [wrapper](#) for Fortran code.

# f2py — combining Fortran and Python

## Basic idea:

`fortrancode.f90` contains a function or subroutine, e.g.  
function `f1(x)` that returns a single value.

```
$ f2py -m mymodule -c fortrancode.f90
```

This creates a binary file `mymodule.so` that can be used as a Python module.

```
>>> from mymodule import f1  
>>> y = f1(3.)
```

## f2py — function example

/codes/f2py/fcn1.f90

```
function f1(x)
    real(kind=8), intent(in) :: x
    real(kind=8) :: f1
    f1 = exp(x)
end function f1
```

Then we can do...

```
$ f2py -m fcn1 -c fcn1.f90
$ python
>>> import fcn1
>>> fcn1.f1(1.)
2.7182818284590451
```

## f2py — subroutine example

/codes/f2py/sub1.f90

```
subroutine mysub(a,b,c,d)
  real (kind=8), intent(in) :: a,b
  real (kind=8), intent(out) :: c,d
  c = a+b
  d = a-b
end subroutine mysub
```

Then we can do...

```
$ f2py -m sub1 -c sub1.f90
$ python
>>> import sub1
>>> y = sub1.mysub(3., 5.)
>>> print y
(8.0, -2.0)
```

**Note:** Tuple (c, d) is returned by the Python function.

## f2py — Jacobi iteration

`/codes/f2py/jacobi1.f90`

```
subroutine iterate(u0, iters, f, u, n)
```

Takes input array `u0` of length `n` and right hand side array `f` and produces `u` by taking `iters` iterations of Jacobi.

`/codes/f2py/plot_jacobi_iterates.py`

```
# Set u = initial guess;  f = rhs
for nn in range(nplots):
    u = jacobi1.iterate(u, iters_per_plot, f)
    plt.plot(x, u, 'o-')
    plt.draw()
    time.sleep(.5)
```

## Other wrappers...

- **Cython**: Allows writing C code embedded in Python.  
<http://www.cython.org/>
- **Jython**: For Java.  
<http://www.jython.org/>
- **swig**: Connects C and C++ to many other languages  
<http://www.swig.org/>

# Mathematical Software

It is best to **use high-quality software** as much as possible, for several reasons:

- It will take **less time** to figure out how to use the software than to write your own version. (Assuming it's well documented!)
- Good general software has been **extensively tested** on a wide variety of problems.
- Often general software is much **more sophisticated** than what you might write yourself, for example it may provide error estimates automatically, or it may be **optimized** to run fast.



# Software sources

- Netlib: <http://www.netlib.org>
- NIST Guide to Available Mathematical Software:  
<http://gams.nist.gov/>
- Trilinos: <http://trilinos.sandia.gov/>
- DOE ACTS: <http://acts.nersc.gov/>
- PETSc nonlinear solvers:  
<http://www.mcs.anl.gov/petsc/petsc-as/>
- Many others!

Many routines for linear algebra  
using non-iterative methods.

Typical name: XYYZZZ

X is precision

YY is type of matrix, e.g. GE (general), BD (bidiagonal),

ZZZ is type of operation, e.g. SV (solve system),

EV (eigenvalues, vectors), SVD (singular values, vectors)

Many routines for linear algebra.

Typical name: XYYZZZ

X is precision

YY is type of matrix, e.g. GE (general), BD (bidiagonal),

ZZZ is type of operation, e.g. SV (solve system),

EV (eigenvalues, vectors), SVD (singular values, vectors)

## Examples:

DGESV can be used to solve a general  $n \times n$  linear system in double precision.

DGTSV can be used to solve a general  $n \times n$  **tridiagonal** linear system in double precision.

# Installing LAPACK

On Virtual Machine or other Debian or Ubuntu Linux:

```
$ sudo apt-get install liblapack-dev
```

This will include BLAS (but not optimized for your system).

Alternatively can download tar files and compile.

See complete documentation at

<http://www.netlib.org/lapack/>

# The BLAS

## Basic Linear Algebra Subroutines

Core routines used by LAPACK (Linear Algebra Package) and elsewhere.

Generally optimized for particular machine architectures, cache hierarchy.

Can create optimized BLAS using ATLAS (Automatically Tuned Linear Algebra Software)

See notes and <http://www.netlib.org/blas/faq.html>

- Level 1: Scalar and vector operations
- Level 2: Matrix-vector operations
- Level 3: Matrix-matrix operations

# The BLAS

Subroutine names start with:

- S: single precision
- D: double precision
- C: single precision complex
- Z: double precision complex

Examples:

- SAXPY: single precision replacement of  $y$  by  $ax + y$ .
- DDOT: dot product of two vectors
- DGEMV: matrix-vector multiply, general matrices
- DGEMM: matrix-matrix multiply, general matrices
- DSYMM: matrix-matrix multiply, symmetric matrices

# Using libraries

If `program.f90` uses BLAS routines...

```
$ gfortran -c program.f90
```

```
$ gfortran program.o -lblas
```

or can combine as

```
$ gfortran program.f90 -lblas
```

When linking together `.o` files, will look for a file called `libblas.a` (probably in `/usr/lib`).

This is a archived static library.

# Using libraries

If `program.f90` uses BLAS routines...

```
$ gfortran -c program.f90
```

```
$ gfortran program.o -lblas
```

or can combine as

```
$ gfortran program.f90 -lblas
```

When linking together `.o` files, will look for a file called `libblas.a` (probably in `/usr/lib`).

This is a archived static library.

Can specify different library location using  
`-L/path/to/library`.



# Making blas library

Download <http://www.netlib.org/blas/blas.tgz>.

Put this in desired location, e.g. `$HOME/lapack/blas.tgz`

```
$ cd $HOME/lapack
$ tar -zxf blas.tgz      # creates BLAS subdirect
$ cd BLAS
$ gfortran -O3 -c *.f
$ ar cr libblas.a *.o    # creates libblas.a
```

To use this library:

```
$ gfortran program.f90 -lblas \
    -L$HOME/lapack/BLAS
```

Note: Non-optimized Fortran 77 versions.

Better approach would be to use [ATLAS](#).

# Creating LAPACK library

Can be done from source at

<http://www.netlib.org/lapack/>

but somewhat more difficult.

Individual routines and dependencies can be obtained from netlib, e.g. the double precision versions from:

<http://www.netlib.org/lapack/double>

Download .tgz file and untar into directory where you want to use them, or make a library of just these files.

# Memory management for arrays

Often a program needs to be written to handle arrays whose size is not known until the program is running.

Fortran 77 approaches:

- Allocate arrays large enough for any application,
- Use “work arrays” that are partitioned into pieces.

We will look at some examples from LAPACK since you will probably see this in other software!

# Memory management for arrays

Often a program needs to be written to handle arrays whose size is not known until the program is running.

Fortran 77 approaches:

- Allocate arrays large enough for any application,
- Use “work arrays” that are partitioned into pieces.

We will look at some examples from LAPACK since you will probably see this in other software!

The good news:

Fortran 90 allows dynamic memory allocation.

# DGESV — Solves a general linear system

<http://www.netlib.org/lapack/double/dgesv.f>

```
SUBROUTINE DGESV( N, NRHS, A, LDA, IPIV,  
&                B, LDB, INFO )
```

$N$  = size of system (square  $N \times N$ )

$A$  = matrix on input, L,U factors on output,  
dimension(LDA,K) with LDA,  $K \geq N$

LDA = **leading dimension of A**  
(number of rows in declaration of A)

**Example:**

```
real(kind=8) dimension(100,500) :: a  
! fill a(1:20, 1:20) with 20x20 matrix  
n = 20  
lda = 100
```

# DGESV — Solves a general linear system

## Example:

```
real(kind=8), dimension(100,500) :: a
real(kind=8), dimension(200,400) :: b
integer, dimension(600) :: ipiv
! fill a(1:20, 1:20) with 20x20 matrix
! b(1:20, 1:3) with 3 right hand sides

n = 20; nrhs = 3; lda = 100; ldb = 200

call dgesv(n, nrhs, a, lda, ipiv, b, ldb, info)
```

What is passed to `dgesv` is `start_address`, the address of first element of `a`. (Matrix is stored by columns)

Whenever `a(i,j)` appears in code, address is:

$$\text{address} = \text{start\_address} + (j-1)*\text{lda} + (i-1)$$

# DGESV — Solves a general linear system

```
SUBROUTINE DGESV( N, NRHS, A, LDA, IPIV,  
&                B, LDB, INFO )
```

NRHS = number of right hand sides

B = matrix whose columns are right hand side(s) on input  
solution vector(s) on output.

LDB = leading dimension of B.

INFO = integer returning 0 if successful.

A = matrix on input, L,U factors on output,

IPIV = Returns pivot vector (permutation of rows)  
integer, dimension(N)  
Row I was interchanged with row IPIV(I).

# Gaussian elimination as factorization

If  $A$  is nonsingular it can be factored as

$$PA = LU$$

where

$P$  is a permutation matrix (rows of identity permuted),

$L$  is lower triangular with 1's on diagonal,

$U$  is upper triangular.

After returning from `dgesv`:

$A$  contains  $L$  and  $U$  (without the diagonal of  $L$ ),

$IPIV$  gives ordering of rows in  $P$ .



# Gaussian elimination as factorization

Example:

$$A = \begin{bmatrix} 2 & 1 & 3 \\ 4 & 3 & 6 \\ 2 & 3 & 4 \end{bmatrix}$$

$$\begin{bmatrix} 0 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 0 & 0 \end{bmatrix} \begin{bmatrix} 2 & 1 & 3 \\ 4 & 3 & 6 \\ 2 & 3 & 4 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 1/2 & 1 & 0 \\ 1/2 & -1/3 & 1 \end{bmatrix} \begin{bmatrix} 4 & 3 & 6 \\ 0 & 1.5 & 1 \\ 0 & 0 & 1/3 \end{bmatrix}$$

IPIV = (2,3,1)

and A ends up as

$$\begin{bmatrix} 4 & 3 & 6 \\ 1/2 & 1.5 & 1 \\ 1/2 & -1/3 & 1/3 \end{bmatrix}$$

## dgesv examples

See `/codes/lapack/random`.

Sample codes that solve the linear system  $Ax = b$  with a random  $n \times n$  matrix  $A$ , where the value  $n$  is run-time input.

`randomsys1.f90` is with static array allocation.

`randomsys2.f90` is with dynamic array allocation.

## dgesv examples

See `/codes/lapack/random`.

Sample codes that solve the linear system  $Ax = b$  with a random  $n \times n$  matrix  $A$ , where the value  $n$  is run-time input.

`randomsys1.f90` is with static array allocation.

`randomsys2.f90` is with dynamic array allocation.

`randomsys3.f90` also estimates **condition number** of  $A$ .

$$\kappa(A) = \|A\| \|A^{-1}\|$$

Can bound relative error in solution in terms of relative error in data using this:

$$Ax^* = b^* \text{ and } A\tilde{x} = \tilde{b} \Rightarrow \frac{\|\tilde{x} - x^*\|}{\|x^*\|} \leq \kappa(A) \frac{\|\tilde{b} - b^*\|}{\|b^*\|}$$

# Just-in-time compilers for Python

Standard implementation of Python as interpreted language.

Importing `mymodule.py` creates `mymodule.pyc`, which is Bytecode (portable code or pcode):

- One-byte operators with operands,  
Interpreted by software at runtime.

Runs much slower than compiled code that is machine-specific instructions.

# Just-in-time compilers for Python

Standard implementation of Python as interpreted language.

Importing `mymodule.py` creates `mymodule.pyc`, which is Bytecode (portable code or pcode):

- One-byte operators with operands,  
Interpreted by software at runtime.

Runs much slower than compiled code that is machine-specific instructions.

Just-in -time (JIT) compilation: Converts bytecode at runtime into native machine code.

Can sometimes run faster than pre-compiled code.

# Just-in-time compilers for Python

## Examples:

- [PyPy](#) — alternative implementation of Python
- [numba](#) — compiles decorated code to [LLVM](#) (formerly Low Level Virtual Machine, compiler infrastructure)

Included in the [Anaconda Python distribution](#)

# Numba — autojit decorator

```
In [1]: def loopsum(n):  
        x = 0  
        for i in range(n):  
            x = x + i
```

```
In [2]: %timeit loopsum(10000)
```

1000 loops, best of 3: 495 us per loop

# Numba — autojit decorator

```
In [1]: def loopsum(n):  
        x = 0  
        for i in range(n):  
            x = x + i
```

```
In [2]: %timeit loopsum(10000)
```

1000 loops, best of 3: 495 us per loop

```
In [3]: from numba import autojit
```

```
In [4]: @autojit  
def loopsum2(n):  
    x = 0  
    for i in range(n):  
        x = x + i
```

```
In [5]: %timeit loopsum2(10000)
```

1000000 loops, best of 3: 1.5 us per loop



# ASCII vs. binary output

Often need to write out a large array of floats with full precision.

For example, one solution value on 3d grid ...

```
do i=1,n
  do j=1,n
    do k=1,n
      write(21,'(e24.16)') u(i,j,k)
    enddo; enddo; enddo
```

How much disk space does this take?

## ASCII vs. binary output

Often need to write out a large array of floats with full precision.

For example, one solution value on 3d grid ...

```
do i=1,n
  do j=1,n
    do k=1,n
      write(21,'(e24.16)') u(i,j,k)
    enddo; enddo; enddo
```

How much disk space does this take?

A single number such as 0.400000000000000000E+01  
has 24 ASCII characters  $\Rightarrow$  24 bytes per value.

Total  $24n^3$  bytes. E.g.  $100 \times 100 \times 100$  grid:  $n = 100 \Rightarrow 24 \text{ MB}$ .

# ASCII vs. binary output

Often need to write out a large array of floats with full precision.

For example, one solution value on 3d grid ...

```
do i=1,n
  do j=1,n
    do k=1,n
      write(21,'(e24.16)') u(i,j,k)
    enddo; enddo; enddo
```

How much disk space does this take?

A single number such as 0.400000000000000000E+01  
has 24 ASCII characters  $\Rightarrow$  24 bytes per value.

Total  $24n^3$  bytes. E.g.  $100 \times 100 \times 100$  grid:  $n = 100 \Rightarrow 24 \text{ MB}$ .

**Note:** In memory storing one 8-byte float takes only 8 bytes.

$(n = 100 \Rightarrow 8\text{MB.})$     **ASCII takes 3 $\times$  the space.**

# ASCII vs. binary output

Often need to write out a large array of floats with full precision.

For example, one solution value on 3d grid ...

```
do i=1,n
  do j=1,n
    do k=1,n
      write(21,'(e24.16)') u(i,j,k)
    enddo; enddo; enddo
```

How much disk space does this take?

A single number such as 0.400000000000000000E+01  
has 24 ASCII characters  $\Rightarrow$  24 bytes per value.

Total  $24n^3$  bytes. E.g.  $100 \times 100 \times 100$  grid:  $n = 100 \Rightarrow 24 \text{ MB}$ .

**Note:** In memory storing one 8-byte float takes only 8 bytes.

$(n = 100 \Rightarrow 8 \text{ MB.})$  **ASCII takes 3 $\times$  the space.**

Also takes additional time to convert to ASCII,

**$\approx 10\times$  slower to write ASCII than dumping binary.**

# Binary output in Fortran

Can use **unformatted** write in Fortran:

```
! /codes/io/binwrite.f90  
open(unit=20, file="u.bin", form="unformatted", &  
      status="unknown", access="stream")  
do j=1,100  
  do i=1,500  
    u(i,j) = real(m*(j-1) + i, kind=8)  
  enddo  
enddo  
  
write(20) u    ! writes entire array in binary  
close(20)
```

-----  
Ascii is 3x larger in size

The resulting binary file **u.bin** cannot be edited directly. But we can read it into Python...

# Reading binary data files in Python

To recover  $U$  array of dimension  $m \times n$  in Python:

```
# /codes/io/binread.py

import numpy as np

file = open('u.bin', 'rb')
uvec = np.fromfile(file, dtype=np.float64)

m,n = np.loadtxt('mn.txt',dtype=int)

# now use Fortran ordering to fill u by columns:
u = uvec.reshape((m,n),order='F')
```

## Other options for binary data

Binary formats that contain a lot of [metadata](#)...

Hierarchical Data Format: HDF, HDF4, HDF5

HDF5 file structure includes two major types of object:

- [Datasets](#): multidimensional arrays of a homogenous type
- [Groups](#): container structures for datasets and other groups

See also: [h5py](#), [PyTables](#)

## Other options for binary data

Binary formats that contain a lot of [metadata](#)...

Hierarchical Data Format: HDF, HDF4, HDF5

HDF5 file structure includes two major types of object:

- [Datasets](#): multidimensional arrays of a homogenous type
- [Groups](#): container structures for datasets and other groups

See also: [h5py](#), [PyTables](#)

NetCDF (Network Common Data Form): Built on top of HDF5.

See also [ncdump](#), [netcdf4-python](#)



## Summary, take away messages...

- Version control — git

Use for all your projects, collaborations, ...

Consider contributing to open source projects

Submit a pull request

## Summary, take away messages...

- **Version control — git**  
Use for all your projects, collaborations, ...  
Consider contributing to open source projects  
Submit a pull request
- **Python, NumPy, SciPy, matplotlib, IPython**  
Quickly trying out new ideas, optimize later  
Graphics and visualization  
Scripting to guide big computations  
Combining codes from different languages  
Many capabilities not seen in class, e.g.  
Manipulating text files, regular expressions,  
building web interfaces

## Summary, take away messages...

- Fortran 90

Compiled language

Tightly constrained but can run very fast

Native multi-dimensional arrays

## Summary, take away messages...

- Fortran 90

Compiled language

Tightly constrained but can run very fast

Native multi-dimensional arrays

- Makefiles

Dependency checking

Often used for building software

# Summary, take away messages...

- Fortran 90

Compiled language

Tightly constrained but can run very fast

Native multi-dimensional arrays

- Makefiles

Dependency checking

Often used for building software

- Debugging code

Unit tests, pytest

Print statements, pdb, gdb

## Summary, take away messages...

- Fortran 90

Compiled language

Tightly constrained but can run very fast

Native multi-dimensional arrays

- Makefiles

Dependency checking

Often used for building software

- Debugging code

Unit tests, pytest

Print statements, pdb, gdb

- Memory hierarchy, cache considerations

Consider layout of arrays in memory

Aim for spatial and temporal locality

# Summary, take away messages...

- **Parallel computing**

Increasingly necessary for all computing

Amdahl's law —

inherently sequential code limits parallelization

Weak vs. strong scaling

Fine grain vs. coarse grain parallelism

Load balancing

# Summary, take away messages...

- **Parallel computing**

Increasingly necessary for all computing

Amdahl's law —

inherently sequential code limits parallelization

Weak vs. strong scaling

Fine grain vs. coarse grain parallelism

Load balancing

- **OpenMP**

Assumes shared memory

Often very easy to add to existing codes

Need to worry about shared/private variables,  
race conditions



## Summary, take away messages...

- MPI — Message Passing Interface

Always assumes distributed memory

Sharing data requires message passing

SPMD: Single Program Multiple Data

Entire program run by each process

But different processes may take different branches

## Summary, take away messages...

- **MPI — Message Passing Interface**

Always assumes distributed memory

Sharing data requires message passing

SPMD: Single Program Multiple Data

Entire program run by each process

But different processes may take different branches

- **Computer arithmetic**

Floating point number representation, 4 byte vs. 8 byte

IEEE standards

Reproducibility still difficult in parallel

Relative error and precision possible

## Summary, take away messages...

- Linear algebra

LAPACK, BLAS — optimized code

Iterative methods for large sparse system

Poisson problems:  $u_{xx} = f(x)$

Two-dimensional Poisson problem  $u_{xx} + u_{yy} = f(x, y)$

## Summary, take away messages...

- Linear algebra

LAPACK, BLAS — optimized code

Iterative methods for large sparse system

Poisson problems:  $u_{xx} = f(x) \Rightarrow$  tridiagonal

Two-dimensional Poisson problem  $u_{xx} + u_{yy} = f(x, y)$

- Quadrature methods / numerical integration

Midpoint, Trapezoid, Simpson Rules

Monte Carlo methods in high dimensions

## Summary, take away messages...

- Linear algebra

Matrix norms and condition number of  $Ax = b$

LAPACK, BLAS — optimized code

Iterative methods for large sparse system

Poisson problems:  $u_{xx} = f(x) \Rightarrow$  tridiagonal

Two-dimensional Poisson problem  $u_{xx} + u_{yy} = f(x, y)$

- Quadrature methods / numerical integration

Midpoint, Trapezoid, Simpson Rules

Monte Carlo methods in high dimensions

- Monte Carlo methods

Pseudo Random Number Generation

Use of seed for reproducibility

Random walks

Happy Computing!